

HW-2: CNN实现CIFAR-10图像分类

姓名：罗威 学号：SA24218095


2.1 整体定义及结果总结

本次作业旨在使用ResNet网络对CIFAR-10数据集进行图像分类，通过调整网络结构、训练参数和优化策略，提高模型的分类准确率，并观察学习曲线的变化。

代码文件： `resnet.py`

训练测试时的终端输出： `terminal_output.txt`

训练的9层ResNet的测试结果表明，该模型在测试集上的正确率达到 **91.89%**。可视化曲线如下图：



```
1 Average Test Loss: 0.3073
2 Accuracy on the 10000 test images: 91.89 %
3
4 --- Per-class Accuracy ---
5 Accuracy of plane : 92.70 %
6 Accuracy of car : 95.80 %
7 Accuracy of bird : 88.60 %
8 Accuracy of cat : 85.00 %
9 Accuracy of deer : 91.70 %
10 Accuracy of dog : 87.40 %
11 Accuracy of frog : 94.20 %
12 Accuracy of horse : 93.10 %
13 Accuracy of ship : 95.90 %
14 Accuracy of truck : 94.50 %
15
16 --- Per-class Precision ---
17 Precision of plane : 92.51 % (TP: 927, Predicted as plane: 1002)
18 Precision of car : 96.28 % (TP: 958, Predicted as car: 995)
19 Precision of bird : 90.22 % (TP: 886, Predicted as bird: 982)
20 Precision of cat : 84.08 % (TP: 850, Predicted as cat: 1011)
21 Precision of deer : 90.79 % (TP: 917, Predicted as deer: 1010)
22 Precision of dog : 87.84 % (TP: 874, Predicted as dog: 995)
23 Precision of frog : 93.27 % (TP: 942, Predicted as frog: 1010)
24 Precision of horse : 95.49 % (TP: 931, Predicted as horse: 975)
25 Precision of ship : 96.09 % (TP: 959, Predicted as ship: 998)
26 Precision of truck : 92.47 % (TP: 945, Predicted as truck: 1022)
```

2.2 实验环境

- **编程语言:** Python
- **PyTorch:** 2.5.0+cu124
- **数据集:** CIFAR-10
- **硬件设备:** GPU: RTX 4060

2.3 训练步骤

2.3.1 库版本检查

检查PyTorch和Torchvision的版本，以及CUDA是否可用。

```
1 print(f"PyTorch Version: {torch.__version__}")
2 print(f"Torchvision Version: {torchvision.__version__}")
3 print(f"CUDA Available: {torch.cuda.is_available()}")
```

2.3.2 使用GPU并行训练

根据CUDA的可用性选择使用GPU或CPU进行训练。

```
1 DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

2.3.3 加载和预处理数据

- 定义训练集和测试集的预处理步骤，包括随机裁剪、随机水平翻转、转换为张量和归一化。
- 下载CIFAR-10数据集，并创建数据加载器。

[illegible]

```

24 classes = ('plane', 'car', 'bird', 'cat', 'deer',
25            'dog', 'frog', 'horse', 'ship', 'truck')
26

```

2.3.4 定义带残差连接的DeepCNN：ResNet

- 定义BasicBlock模块，包含两个卷积层和一个残差连接。
- 定义ResNet网络，由多个BasicBlock模块组成，并包含全局平均池化和全连接层。
- 定义ResNet18和ResNet9两种网络结构。

```

1  class BasicBlock(nn.Module):
2      expansion = 1 # BasicBlock的输出通道数与目标输出通道数相同
3      def __init__(self, in_planes, planes, stride=1):
4          super(BasicBlock, self).__init__()
5          self.conv1 = nn.Conv2d(in_planes, planes, kernel_size=3,
stride=stride, padding=1, bias=False)
6          self.bn1 = nn.BatchNorm2d(planes)
7          self.conv2 = nn.Conv2d(planes, planes, kernel_size=3, stride=1,
padding=1, bias=False)
8          self.bn2 = nn.BatchNorm2d(planes)
9
10         self.shortcut = nn.Sequential()
11         # 如果输入输出通道数不同，或者stride不为1（特征图尺寸改变），就使用一层卷积和
BatchNorm来修正维度
12         if stride != 1 or in_planes != self.expansion*planes:
13             self.shortcut = nn.Sequential(
14                 nn.Conv2d(in_planes, self.expansion*planes, kernel_size=1,
stride=stride, bias=False),
15                 nn.BatchNorm2d(self.expansion*planes)
16             )
17         def forward(self, x):
18             out = F.relu(self.bn1(self.conv1(x)))
19             out = self.bn2(self.conv2(out))
20             out += self.shortcut(x) # 残差连接
21             out = F.relu(out)
22             return out
23
24  class ResNet(nn.Module):
25      def __init__(self, block, num_blocks, num_classes=10):
26          super(ResNet, self).__init__()
27          self.in_planes = 64 # 初始通道数
28
29          # 初始卷积层：3x3的卷积
30          self.conv1 = nn.Conv2d(3, self.in_planes, kernel_size=3, stride=1,
padding=1, bias=False)
31          self.bn1 = nn.BatchNorm2d(self.in_planes)
32
33          self.layer1 = self._make_layer(block, 64, num_blocks[0], stride=1)
34          self.layer2 = self._make_layer(block, 128, num_blocks[1], stride=2)
# stride=2 实现下采样
35          self.layer3 = self._make_layer(block, 256, num_blocks[2], stride=2)
36          self.layer4 = self._make_layer(block, 512, num_blocks[3], stride=2)
37
38          self.avgpool = nn.AdaptiveAvgPool2d((1, 1)) # 全局平均池化
39          self.linear = nn.Linear(512*block.expansion, num_classes)

```

```

40
41     def _make_layer(self, block, planes, num_blocks, stride):
42         strides = [stride] + [1]*(num_blocks-1)
43         layers = []
44         for s in strides:
45             layers.append(block(self.in_planes, planes, s))
46             self.in_planes = planes * block.expansion
47         return nn.Sequential(*layers)
48
49     def forward(self, x):
50         out = F.relu(self.bn1(self.conv1(x)))
51         out = self.layer1(out)
52         out = self.layer2(out)
53         out = self.layer3(out)
54         out = self.layer4(out)
55         out = self.avgpool(out)
56         out = torch.flatten(out, 1)
57         out = self.linear(out)
58         return out
59
60     # 4. 初始化ResNet配置
61     # 4.1 这是一个简化的ResNet18结构, 适应CIFAR-10
62     def ResNet18():
63         # ResNet18的block配置是 [2, 2, 2, 2]
64         return ResNet(BasicBlock, [2, 2, 2, 2], num_classes=10)
65
66     # 4.2 一个更小的ResNet变体, 包括全连接层在内一共9层
67     def ResNet9():
68         return ResNet(BasicBlock, [1,1,1,1], num_classes=10)

```

2.3.5 初始化ResNet配置

- 选择ResNet9作为训练模型, 并将其移动到指定的设备上。

```

1 # model = ResNet18().to(DEVICE)
2
3 # 使用较小的ResNet来进行训练, 因为这个网络的性能已经足够
4 model = ResNet9().to(DEVICE)
5 print(model)

```

2.3.6 定义损失函数和优化器

- 使用交叉熵损失函数作为损失函数。
- 使用随机梯度下降 (SGD) 优化器, 并设置学习率、动量和权重衰减。
- 使用余弦退火调度器调整学习率。

```

1 criterion = nn.CrossEntropyLoss()
2 # 优化器: 对于ResNet结构, SGD with momentum比Adam优化器更好
3 optimizer = optim.SGD(model.parameters(), lr=LEARNING_RATE, momentum=0.9,
4 weight_decay=WEIGHT_DECAY)
5
6 # 学习率调度器: CosineAnnealingLR--余弦退火调度来调整学习率
7 scheduler = optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=EPOCHS)

```

2.3.7 训练模型

- 定义训练函数，包括前向传播、计算损失、反向传播和更新参数。
- 在训练过程中，记录训练损失和准确率，并打印每个epoch的训练结果。

```
1 def train_model(model, trainloader, criterion, optimizer, scheduler,
2 epochs):
3     print("\n--- Training Started ---")
4     train_losses = []
5     train_accuracies = []
6
7     for epoch in range(epochs):
8         model.train()
9         running_loss = 0.0
10        correct_train = 0
11        total_train = 0
12
13        current_lr = optimizer.param_groups[0]['lr']
14        print(f"Epoch [{epoch+1}/{epochs}], Current Learning Rate:
15        {current_lr:.6f}")
16
17        for i, data in enumerate(trainloader, 0):
18            inputs, labels = data
19            inputs, labels = inputs.to(DEVICE), labels.to(DEVICE)
20
21            optimizer.zero_grad()
22            outputs = model(inputs)
23            # 1. 计算loss
24            loss = criterion(outputs, labels)
25            loss.backward()
26            optimizer.step()
27
28            # 2. 累加每个batch的总损失
29            running_loss += loss.item() * inputs.size(0)
30
31            # 3. 计算准确率
32            _, predicted = torch.max(outputs.data, 1)
33            total_train += labels.size(0)
34            correct_train += (predicted == labels).sum().item()
35
36            # 4. 正确计算epoch平均损失
37            epoch_train_loss = running_loss / total_train
38            epoch_train_acc = 100 * correct_train / total_train
39
40            train_losses.append(epoch_train_loss)
41            train_accuracies.append(epoch_train_acc)
42
43            print(f"Epoch {epoch+1} finished. Avg Training Loss:
44            {epoch_train_loss:.4f}, Training Accuracy: {epoch_train_acc:.2f}%")
45            scheduler.step()
46
47        return train_losses, train_accuracies
48
49 train_losses, train_accuracies = train_model(model, trainloader, criterion,
50 optimizer, scheduler, EPOCHS)
```

2.3.8 测试模型并计算评价指标

- 定义测试函数，包括前向传播、计算损失和准确率。
- 在测试过程中，计算每一类别的准确率和精度，并打印测试结果。

```
1  # 测试模型并计算评价指标
2  def test_model_and_evaluate(model, testloader, classes):
3      print("\n--- Testing Started ---")
4      model.eval()
5      correct = 0
6      total = 0
7      test_loss = 0.0
8
9      class_true_positives = defaultdict(int)
10     class_predicted_positives = defaultdict(int)
11     class_correct = list(0. for i in range(len(classes)))
12     class_total = list(0. for i in range(len(classes)))
13
14
15     with torch.no_grad():
16         for data in testloader:
17             images, labels = data
18             images, labels = images.to(DEVICE), labels.to(DEVICE)
19             outputs = model(images)
20             loss = criterion(outputs, labels)
21             test_loss += loss.item()
22
23             _, predicted = torch.max(outputs.data, 1)
24             total += labels.size(0)
25             correct += (predicted == labels).sum().item()
26
27         # 计算每一类别的准确率
28         c = (predicted == labels).squeeze()
29         for i in range(labels.size(0)):
30             label = labels[i]
31             class_correct[label] += c[i].item()
32             class_total[label] += 1
33
34         # 计算每一类别的精度
35         for i in range(labels.size(0)):
36             label = labels[i].item()
37             pred = predicted[i].item()
38             class_predicted_positives[pred] += 1
39             if label == pred:
40                 class_true_positives[label] += 1
41
42     avg_test_loss = test_loss / len(testloader)
43     overall_accuracy = 100 * correct / total
44     print(f'\nAverage Test Loss: {avg_test_loss:.4f}')
45     print(f'Accuracy on the {total} test images: {overall_accuracy:.2f} %')
46
47     print("\n--- Per-class Accuracy ---")
48     for i in range(len(classes)):
49         if class_total[i] > 0:
```

```

50         print(f'Accuracy of {classes[i]:5s} : {100 * class_correct[i] /
class_total[i]:.2f} %')
51     else:
52         print(f'Accuracy of {classes[i]:5s} : N/A (no instances in test
set)')
53
54
55     print("\n--- Per-class Precision ---")
56     for i in range(len(classes)):
57         class_name = classes[i]
58         tp = class_true_positives[i]
59         tp_plus_fp = class_predicted_positives[i]
60         precision = 0
61         if tp_plus_fp > 0:
62             precision = 100 * tp / tp_plus_fp
63         print(f'Precision of {class_name:5s} : {precision:.2f} % (TP: {tp},
Predicted as {class_name}: {tp_plus_fp})')
64
65     print('--- Finished Testing ---')
66     return avg_test_loss, overall_accuracy
67
68 avg_test_loss, test_accuracy = test_model_and_evaluate(model, testloader,
classes)
69

```

2.3.9 绘制学习曲线

- 绘制训练损失和准确率随epoch的变化曲线，并保存为图像文件。

```

1  matplotlib.use('Agg')
2  plt.figure(figsize=(12, 4))
3  plt.subplot(1, 2, 1)
4  plt.plot(range(1, EPOCHS + 1), train_losses, label='Training Loss')
5  plt.title('Training Loss over Epochs')
6  plt.xlabel('Epochs')
7  plt.ylabel('Loss')
8  plt.legend()
9
10 plt.subplot(1, 2, 2)
11 plt.plot(range(1, EPOCHS + 1), train_accuracies, label='Training Accuracy')
12 plt.title('Training Accuracy over Epochs')
13 plt.xlabel('Epochs')
14 plt.ylabel('Accuracy (%)')
15 plt.legend()
16 plt.tight_layout()
17 plt.tight_layout()
18 plt.savefig('result.png', # 保存路径
19             dpi=300,      # 分辨率
20             bbox_inches='tight', # 去除多余白边
21             format='png'    # 输出格式
22             )

```

2.4 训练结果

2.4.1 模型训练结果

- 训练损失随epoch的增加逐渐减小，最终收敛到较低水平。
- 训练准确率随epoch的增加逐渐提高，最终达到**99.46%**。

2.4.2 模型测试结果

- 测试损失和准确率与训练结果相似，表明模型具有较好的泛化能力。
- 每一类别的准确率和精度最低为**85%**，最高为**95.9%**，表现出很好的识别能力。

2.4.3 学习曲线

- 训练损失曲线和准确率曲线呈现出典型的收敛趋势，表明模型的训练过程是稳定的。

